One primary reason I am excited about academia is the unique opportunity it offers to influence the next generation of researchers through teaching and mentorship. I am grateful for having received excellent guidance from all my mentors and teachers, which has, in no small part, shaped me and inspired me to give back to the community. Building on my own journey over the last several years as a teacher and mentor, I aspire to contribute in an even larger capacity in the future.

Teaching Experience My experiences with teaching began while I was a PhD student at Carnegie Mellon University. In the Spring of 2015, I was the lead teaching assistant for the CMU LTI 11-711 Natural Language Processing course, with a large attendance of over 70 undergraduate and graduate students. As part of my responsibilities, I held recitations for lectures, and regular discussion hours following the class, graded homework assignments and exams. While the strength of the class made my responsibilities challenging at times, it was nevertheless exciting to see what happens behind the curtains—how lectures and assignments are created and graded, and most of all how intuitions about complex topics are conveyed to students. One of the most gratifying experiences during recitations was receiving questions that revealed the student's thought process, based on which I would modify my answers catering to their learning needs.

I have also been fortunate to be invited as a guest lecturer for various NLP and ML classes, across different universities, listed below:

- Dependency Parsing with Chu-Liu-Edmonds [Fall 2014, CMU 11-271, *Algorithms for NLP*]
- Unsupervised Learning [Fall 2017, UW CSE 446, *Machine Learning*]
- Dependency Parsing [Spring 2018, UW CSEP 517, *NLP*]
- Minimum Bayes Risk Decoding [Spring 2018, UW CSE 599 D1, *Advanced Topics in NLP*]
- Phrase-Structure Parsing [Winter 2019, UW CSE 447/547, *NLP*]
- Interpretability and Biases in NLP [Fall 2020, UT-Austin CS395T, *Topics in NLP*]
- Interpretability in NLP [Winter 2021 (upcoming), UW CSE 447/547, *NLP*]
- Transfer Learning [Winter 2021 (upcoming), UW DATA 598, *Statistical Deep Learning*]

I have made a conscious effort in all my lectures to engage with students, and am always pleasantly surprised by how enjoyable the interactions turn out for myself. Based on this experience, I am looking forward to an opportunity to design courses in NLP and AI, especially for the stimulating challenge of building a curriculum for such rapidly changing, yet immensely relevant fields.

I have had the honor to reach an even wider audience through tutorials at two different conference venues. I co-taught the tutorial on Frame Semantics Across Languages with researchers from UC Berkeley at COLING 2018. This was primarily a linguistics venue, and I treated it as a compelling opportunity to teach concepts from NLP and machine learning. At NAACL 2019, I co-taught the tutorial on Transfer Learning in NLP. It received an audience of about 350 people, and was one of the most-attended tutorials at the conference. I have also co-organized a workshop on deep learning approaches for low-resource NLP (DeepLo 2019), collocated with EMNLP 2019. All the above experiences have taught me how to convey research fundamentals, through first principles, to a more general audience, who may or may not have familiarity with my research area.

All throughout my research career, I have led reading groups attended by students, postdocs, faculty and researchers, where we discussed conference and journal papers. In my current role at the Allen Institute for AI, I organize and moderate the weekly research meetings for the MOSAIC team. In each of these meetings I try to engage the audience by synthesizing ideas from the presentation and their own broader research interests.

Mentorship I have had the privilege of mentoring many talented undergraduate, Masters and PhD students during my research career. Two of these efforts have even led to theses: Karishma Mandyam wrote her Senior thesis and Ron Fan wrote his Masters thesis, based on their research with me. The collaborations with the graduate students have led to publications at ACL 2018, NAACL 2018, EMNLP 2020, with another under submission at EACL 2020. In my current role as a postdoc at the Allen Institute for AI, I have continued my mentorship efforts. Over summer and fall, I have been mentoring three interns: Jize Cao, Jenny Liang, and Alon Jacovi on various projects related to interpretability in

NLP and AI. I have also mentored two pre-doctoral residents during the last year, leading to publications at EMNLP 2020, and at ACL 2020; the latter also received the honorable mention award at ACL. In each of these situations, I have guided the students by starting from high-level research questions, helped them break these down into concrete experiments, and finally assisted them with technical writing for papers, reports or theses. These experiences have taught me that learning approaches by individual students can be varied; each student has their own unique journey in their growth as scientists. However, it has been uniformly fulfilling to watch my students grow and mature in their scientific pursuits.

I have additionally volunteered as a mentor at conferences, namely EMNLP 2018, ACL 2020 and EMNLP 2020 to help incoming and continuing PhD students from all over the world with questions regarding success in graduate school admissions and completion. With a spirit of giving back to the community that helped me get into graduate school, I have helped upwards of 15 PhD and Masters applicants with their application materials over the past several years.

Teaching Interests Building on my prior teaching experience and research interests, I would be excited to teach courses both at the undergraduate and graduate levels, in any of the following areas: Natural Language Processing, Artificial Intelligence, Machine Learning, Deep Learning, Data Science, and Ethics in Artificial Intelligence. I would also be interested in developing research courses in relatively new research areas in NLP, such as model and data interpretability, either at the graduate level or as project-focused, senior undergraduate seminars. Based on my experience leading reading groups and research discussions at UW and the Allen Institute for AI respectively, I would love to organize seminar classes where I can encourage synthesis of ideas from research papers to inspire research-minded students.

Building on my outreach experience with UW CSE GradsWomen and the Highline Public High School near Seattle, I plan to seek opportunities for community engagement and outreach in computer science, where certain demographic groups tend to be overlooked. As a woman of South Asian origin, and an immigrant in the United States, I firmly believe that every individual has strengths shaped by their unique background, even though they might not be immediately apparent via traditional metrics of evaluation. By providing opportunities to folks from varied backgrounds, we can empower and unlock the potential of a diverse community of researchers who will ask and answer questions pertinent to the broader society.